

1 Definitions

Non-collapsing hash function. A hash function that is collision-resistant and not infinitely-often collapsing. (See [Zha17] for definitions of ‘collision-resistant’ and ‘infinitely-often collapsing’. We can assume that there are adversaries who win the collapsing game for such functions with probability near 1. See [GYZ17], [Zha17] for examples of how one can boost the success probability of any bad adversary who breaks infinitely-often collapsing security so that it becomes a very good adversary.)

Part-probabilistic non-collapsing hash function. A function $F(k, x, r)$ whose output takes the form $H(k, x) \parallel R(H(k, x), r)$, where H is a (deterministic) collision-resistant hash function, and R is another function whose output is permitted to depend upon the randomness r and the output of H . We require the following security properties of F :

1. Collision resistance for H : For any quantum polynomial time adversary A ,

$$\Pr[H(k, x_0) = H(k, x_1) \wedge x_0 \neq x_1 : (x_0, x_1) \leftarrow A(k), k \leftarrow \{0, 1\}^\lambda] < \text{negl}(\lambda)$$

2. (Infinitely-often) non-collapsing: There exists an adversary A (consisting of two phases, A_0 and A_1) who can win the following game with probability $1 - \gamma$, where γ is negligible.

- The challenger has an input bit b .
- The challenger chooses a random key k , which it gives to A_0 .
- A_0 creates a superposition $|\psi\rangle = \sum_x \alpha_x |x\rangle$ and submits this state to the challenger.
- The challenger generates a random r . It evaluates $F(k, \cdot, r)$ in superposition on $|\psi\rangle$, to get the state $\sum_x \alpha_x |x, H(k, x), R(H(k, x), r)\rangle$.
- The challenger does one of the following:
 - If $b = 0$, it measures the last two registers, and returns the state $\sum_{x: H(k, x) = y} \alpha_x |x, y, R(y, r)\rangle$ to A .
 - If $b = 1$, it measures the entire state, and returns the state $|x_0, y, R(y, r)\rangle$ (for some x_0) to A .
- A_1 outputs a guess for b . If A_1 is correct, A wins the game.

Note that, under this definition of ‘non-collapsing’ (which mimics [Zha17]’s definition), the distinguisher A_1 is only guaranteed to exist *if the challenger behaves honestly*. A_1 ’s success may depend upon r being honestly generated, and upon R being honestly run; it has no way of verifying that either is the case. In the generic hash function setting, we cannot guarantee that R will be run honestly on a random r , and we cannot guarantee that the distinguisher A_1 which wins the game above will still be useful if this is not the case.

It is evident that the (deterministic) non-collapsing hash function is a special case of the part-probabilistic non-collapsing hash function, so that any NCH function is also a PP-NCH function.

Chosen- y -secure hash function. A hash function $H(k, x)$ for which no quantum polynomial time adversary can win the following game with more than negligible probability:

- The challenger chooses a random key k , which it gives to A .
- The challenger creates a uniform superposition over all inputs x in the input space of H , and evaluates $H(k, \cdot)$ upon this superposition to obtain the state $\sum_x \alpha_x |x, H(k, x)\rangle$. It then measures the output register to obtain a state $|\psi_y\rangle = \sum_{x: H(k, x) = y} \alpha_x |x, y\rangle$ for some random y .
- The challenger gives $|\psi_y\rangle$ to the adversary. The adversary wins the game if it can recover x_0, x_1 such that $H(k, x_0) = H(k, x_1) = y$.

The collapsing security game can be defined in the same way for CYS hash functions that it is for collision-resistant hash functions.

Note that it is easy to construct a collision-resistance adversary from a chosen- y adversary, and that, therefore, any collision-resistant hash function is also a chosen- y -secure hash function. This notion of ‘chosen y ’ security is close to that of second preimage resistance security; the former can be considered a strengthening of the latter to suit the quantum setting.

Note, in addition, that any PP-NCH function can be transformed into a CYS-NCH function. Chosen- y security follows directly from the assumption of collision resistance for the deterministic part of the PP-NCH function. The non-collapsing property, meanwhile, follows from the fact that the chosen- y setting already demands that y is chosen randomly (presumably by some trusted party) if the CYS function’s preimage security is to hold. Any construction using a chosen- y hash function for its preimage security properties, therefore, must generate honest randomness when it generates y ; and, as such, y can be used as a source of randomness for the randomised part of the PP-NCH function.

2 NCH implies BZ-GYZ

Claim. Any non-collapsing hash function can be used to build a one-time signature scheme that is Boneh-Zhandry secure but not Garg-Yuen-Zhandry secure.

Scheme. Given a non-collapsing hash function F , and an arbitrary BZ-secure one-time signature scheme $(\text{Gen}_{\text{BZ}}, \text{Sign}_{\text{BZ}}, \text{Ver}_{\text{BZ}})$ (these are easy to produce; for example, the standard Lamport construction of a one-time signature scheme from a collision-resistant hash function is BZ-secure, according to [BZ13]), we construct a BZ-GYZ scheme $(\text{Gen}, \text{Sign}, \text{Ver})$ as follows.

- **Gen** simply runs Gen_{BZ} to generate a pair of keys (pk, sk) for the BZ signature scheme.
- $\text{Sign}(\text{sk}, m) = F(m) \parallel \text{Sign}_{\text{BZ}}(\text{sk}, F(m)) = \sigma$. In other words, **Sign** applies the non-collapsing hash function F to the message, signs the hashed message using the BZ scheme, and outputs the hashed message concatenated with the signature it obtains from the BZ signing oracle.
- $\text{Ver}(\text{pk}, m, \sigma)$ firstly hashes the message m to obtain $F(m)$, and then verifies $F(m)$ using Ver_{BZ} and pk .

Proof. We firstly prove that this scheme is BZ-secure, assuming that $(\text{Gen}_{\text{BZ}}, \text{Sign}_{\text{BZ}}, \text{Ver}_{\text{BZ}})$ is BZ-secure.

Suppose we have some adversary A who is able to break the BZ-security of $(\text{Gen}, \text{Sign}, \text{Ver})$. The adversary B can then use A to break the BZ-security of $(\text{Gen}_{\text{BZ}}, \text{Sign}_{\text{BZ}}, \text{Ver}_{\text{BZ}})$ as follows:

- B receives pk from its challenger. It passes pk on to A .
- A creates a superposition of messages $\sum_m \alpha_m |m\rangle$ and gives it to B as a query. B computes F on it in superposition, and then passes $\sum_m \alpha_m |m, F(m)\rangle$ on to its challenger, who computes Sign_{BZ} on it and returns the state $\sum_m \alpha_m |m, F(m), \text{Sign}_{\text{BZ}}(\text{sk}, F(m))\rangle$ to B . B gives this state to A .
- A outputs its (classical) forgery for $(\text{Gen}, \text{Sign}, \text{Ver})$. This forgery will take the form $((m_0, F(m_0), \sigma_0), (m_1, F(m_1), \sigma_1))$, where m_0 and m_1 are two distinct messages.
- If $F(m_0) = F(m_1)$, then we have found a collision for F , which ought to be impossible, because we assume that F , a non-collapsing hash function, is collision-resistant. If $F(m_0) \neq F(m_1)$, then B outputs $((F(m_0), \sigma_0), (F(m_1), \sigma_1))$ as its forgery for the $(\text{Gen}_{\text{BZ}}, \text{Sign}_{\text{BZ}}, \text{Ver}_{\text{BZ}})$ scheme.
- If A ’s success probability is non-negligible, then so is B ’s.

Therefore, $(\text{Gen}, \text{Sign}, \text{Ver})$ is BZ-secure if $(\text{Gen}_{\text{BZ}}, \text{Sign}_{\text{BZ}}, \text{Ver}_{\text{BZ}})$ is BZ-secure.

We now prove that the proposed scheme is not GYZ-secure. To do this, we use the fact that F is non-collapsing. Let D be an adversary which can break collapsing security for F . Following the proof to Theorem 13 in [GYZ17], we construct an adversary A who uses D to break the GYZ-security of $(\text{Gen}, \text{Sign}, \text{Ver})$. A acts as follows:

- A queries D to get a superposition of ‘messages’ (preimages) $\rho = \sum_{m,m'} \alpha_m \alpha_{m'}^* |m\rangle\langle m'|$, and places this superposition ρ in its message register. A then flips a coin with outputs in $\{0, 1\}$ and measures the message register iff the coin gives 0. This results in the following state:

$$\frac{1}{2} \left[|0\rangle\langle 0| \otimes \sum_m |\alpha_m|^2 |m\rangle\langle m| + |1\rangle\langle 1| \otimes \rho \right] \quad (1)$$

- A sends this state to the GYZ signing oracle, which signs it and places the signature in a newly created pair of signature registers. The result is then

$$\frac{1}{2} \left[\begin{aligned} & \sum_m |\alpha_m|^2 \quad |0\rangle\langle 0| \otimes |m\rangle\langle m| \otimes |F(m)\rangle\langle F(m)| \otimes |\text{Sign}_{\text{BZ}}(\text{sk}, F(m))\rangle\langle \text{Sign}_{\text{BZ}}(\text{sk}, F(m))| \\ & + \sum_{m,m'} \alpha_m \alpha_{m'}^* \quad |1\rangle\langle 1| \otimes |m\rangle\langle m'| \otimes |F(m)\rangle\langle F(m')| \otimes |\text{Sign}_{\text{BZ}}(\text{sk}, F(m))\rangle\langle \text{Sign}_{\text{BZ}}(\text{sk}, F(m'))| \end{aligned} \right] \quad (2)$$

A measures the ‘ $F(m)$ ’ register (the third register from the left), to obtain the state

$$\frac{1}{2} \sum_y \beta_y \left[\begin{aligned} & \sum_{m:F(m)=y} |\alpha_m|^2 \quad |0\rangle\langle 0| \otimes |m\rangle\langle m| \otimes |y\rangle\langle y| \otimes |\text{Sign}_{\text{BZ}}(\text{sk}, y)\rangle\langle \text{Sign}_{\text{BZ}}(\text{sk}, y)| \\ & + \quad |1\rangle\langle 1| \otimes |\psi_y\rangle\langle \psi_y| \otimes |y\rangle\langle y| \otimes |\text{Sign}_{\text{BZ}}(\text{sk}, y)\rangle\langle \text{Sign}_{\text{BZ}}(\text{sk}, y)| \end{aligned} \right] \quad (3)$$

where $|\psi_y\rangle = \sum_{m:F(m)=y} \alpha_m |m\rangle$.

- A applies D to the second and third registers from the left and saves D ’s output in a newly created ancilla register. If D is a very good distinguisher which gives the right answer with probability $1 - \gamma$, then, by the gentle measurement lemma, the resulting state is $4\sqrt{2\gamma}$ close to

$$\frac{1}{2} \sum_y \beta_y \left[\begin{aligned} & \sum_{m:F(m)=y} |\alpha_m|^2 \quad |0\rangle\langle 0| \otimes |m\rangle\langle m| \otimes |y\rangle\langle y| \otimes |\text{Sign}_{\text{BZ}}(\text{sk}, y)\rangle\langle \text{Sign}_{\text{BZ}}(\text{sk}, y)| \otimes |0\rangle\langle 0| \\ & + \quad |1\rangle\langle 1| \otimes |\psi_y\rangle\langle \psi_y| \otimes |y\rangle\langle y| \otimes |\text{Sign}_{\text{BZ}}(\text{sk}, y)\rangle\langle \text{Sign}_{\text{BZ}}(\text{sk}, y)| \otimes |1\rangle\langle 1| \end{aligned} \right] \quad (4)$$

(It is possible to boost the success probability of any bad distinguisher so that it becomes a good distinguisher, as long as the bad distinguisher still breaks collapsing security for F . See Section 4.1 of [GYZ17].)

- A sends this state to its challenger for verification. Because all the signatures in this state were legally obtained, this state passes verification with probability 1, and is not perturbed by verification. Therefore, the challenger’s output $\text{GYZ-Exp}(A)$ is equal to the state above.
- By the same reasoning that is used in the proof of Theorem 13 in [GYZ17], there is no basis-respecting adversary that could produce such an output, as basis-respecting adversaries must commute with measurement in the computational basis. (Intuitively, it is clear that the distinguisher does not commute with measurement, because its very purpose is to determine whether or not a state has been measured.)

Therefore, if F is non-collapsing, $(\text{Gen}, \text{Sign}, \text{Ver})$ is not GYZ-secure.

3 NCH implies quantum tokens

Claim. Any non-collapsing hash function can be used to build a testable tokenised signature scheme.

Scheme. Given an arbitrary non-collapsing hash function F , we construct a one-bit, one-time, testable tokenised signature scheme. (By Section 5 of [BS16], a one-bit, one-time tokenised signature scheme can be extended to a fully-fledged tokenised signature scheme via a series of reductions.)

Because F is non-collapsing, there exists an adversary A that can break collapsing security for F with very high success probability $1 - \gamma$. We use A , along with an arbitrary classical digital signature scheme $(\text{Gen}, \text{Sign}, \text{Ver})$, to construct a tokens scheme.

- Let (pk, sk) denote the public and secret keys for the classical signature scheme, and let (pk, sk) stand for the keys to our tokens scheme.
- A consists of two phases: 1) the phase which outputs a superposition of messages for its challenger to hash, and 2) the phase which guesses whether the challenger measured its entire state or only the output registers. Following [Zha17], we let A_0 denote the first phase and A_1 the second.
- **key-gen** runs **Gen** to generate (pk, sk) for the digital signature scheme. It then outputs $sk = (\text{sk}, A_0)$ and $pk = (\text{pk}, A_1)$.
- **token-gen** (sk) firstly runs A_0 , twice, to generate two superpositions of messages $\sum_{m_0} \alpha_{m_0} |m_0\rangle$ and $\sum_{m_1} \alpha_{m_1} |m_1\rangle$. It computes F in superposition on both, and measures both output registers. The result is two states $|\psi_{y_0}\rangle, |\psi_{y_1}\rangle$, where $|\psi_{y_b}\rangle = \sum_{m:F(m)=y_b} |m, F(m)\rangle$. It then signs the tuple (y_0, y_1) using sk and **Sign**. **token-gen** outputs $(|\psi_{y_0}\rangle, |\psi_{y_1}\rangle, \text{Sign}(\text{sk}, (y_0, y_1)))$ as $|\underline{\sigma}\rangle$.
- **sign** $(\alpha \in \{0, 1\}, |\underline{\sigma}\rangle)$ measures the input register of $|\psi_{y_\alpha}\rangle$, and outputs, as a signature for α , a preimage m_α for y_α under F , along with an unaltered $|\psi_{y_{1-\alpha}}\rangle$ and the signature $\text{Sign}(\text{sk}, (y_0, y_1))$.
- **verify** $(pk, \alpha, \sigma = (m_\alpha, |\psi_{y_{1-\alpha}}\rangle, \text{Sign}(\text{sk}, (y_0, y_1))))$ firstly verifies the signature on (y_0, y_1) using pk and **Ver**. Following this, it checks that m_α hashes to y_α , and that the superposition of messages in the message register of $|\psi_{y_{1-\alpha}}\rangle$ hashes to $y_{1-\alpha}$. It then checks, using A_1 , that the purported $|\psi_{y_{1-\alpha}}\rangle$ really does still have an unmeasured message register. In more precise terms, if A_1 outputs ‘measured’, **verify** outputs F ; otherwise, **verify** outputs T .
- **verify-token** $(|\underline{\sigma}\rangle)$ applies **Ver** $_0$, from Section 4.2 of [Zha17], to both candidate states $|\psi_{y_0}\rangle$ and $|\psi_{y_1}\rangle$.

Proof.

Testability. Testability follows directly from Zhandry’s work in Section 4.2 of [Zha17]. The correctness portion of testability is identical to the correctness requirement for quantum money, which Zhandry’s construction satisfies. The security portion of testability also follows from Section 4.2. By Zhandry’s proof of security, we can assume that two dishonest candidate states $|\phi_0\rangle$ and $|\phi_1\rangle$ pass **verify-token** with at most negligible probability. Therefore, except with negligible probability, any two states passing **verify-token** will be honest states. Because **verify** always accepts honest tokens, our scheme satisfies equation (8) of [BS16].

Unforgeability. Suppose that there is an adversary who can, after seeing a single token $|\underline{\sigma}\rangle = (|\psi_{y_0}\rangle, |\psi_{y_1}\rangle)$, produce two signatures $(x_0, |\phi_0\rangle)$ and $(x_1, |\phi_1\rangle)$ that both pass **verify** with non-negligible probability. (Assume, for the present, that the adversary did not attempt to forge a signature on (y_0, y_1) .)

Without loss of generality, consider the $(x_0, |\phi_0\rangle)$ tuple. To pass the hash tests which **verify** executes, x_0 must be a valid preimage to y_0 under F . Note that the only states $|\phi_0\rangle$ which the adversary can hold without violating the collision resistance of F , given that he already has x_0 , are states negligibly close to $|x_0, y_0\rangle$. However, these states will almost invariably fail A_1 ’s distinguishing test, because A_1 is—under the assumption that we have already boosted its success probability to some $1 - \gamma$ —able to tell, with probability $1 - \gamma$, the difference between any $|x_0, y_0\rangle$ and the state $|\psi_{y_0}\rangle$. The adversary’s only alternative is to forge a signature on (y_0, y_1) , but his succeeding would violate the security of $(\text{Gen}, \text{Sign}, \text{Ver})$, which we assume is impossible. Therefore, there is no adversary which can produce two signatures that pass **verify**, having seen only one token—except with negligible probability.

Remarks.

- The part of this scheme which is applicable to public-key quantum money is no different from Zhandry’s construction of [Zha17]. We have essentially produced an extension of that construction which happens to be a tokens scheme.

- One curious property of this NCH-based construction for tokens is that the signatures σ it produces are quantum, while [BS16]’s signatures were classical. In consequence, our construction loses some of the properties which [BS16] considered desirable, such as the ability to convert quantum money into ‘classical cheques’ which could be sent over classical channels. The sacrifice is not a fruitless one, however: since our scheme is based on non-collapsing hash functions, we also acquire the properties of ‘collision-free quantum money’ which [Zha17] considered desirable, such as the ability to ensure that even the bank cannot forge quantum money (or signing tokens).

4 NCH from iO

Let q be a positive even integer. Let L be a map from \mathbb{F}_2^n to \mathbb{Z}_q^n such that

$$L(x \in \mathbb{F}_2^n) = v \in \mathbb{Z}_q^n = \begin{pmatrix} v_1 \\ \vdots \\ v_n \end{pmatrix}, v_i = \begin{cases} 0 & x_i = 0 \\ \frac{q}{2} & x_i = 1 \end{cases}. \quad (5)$$

Note that L is injective.

The non-collapsing hash function is defined as follows:

- **key-gen**: Generates the following constants which define the hash function.

- A randomly-chosen $m \times n$ LWE matrix, A , whose entries are in \mathbb{Z}_q ;
- k n -bit binary strings s_1, \dots, s_k , which represent k secrets.

The hash function key that **key-gen** outputs consists of the following:

- The matrix A .
- k LWE samples $A(L(s_i)) + e_i$, for i in $1, \dots, k$. e_i are randomly chosen error vectors in \mathbb{Z}_q^m with sufficiently small bounded magnitude.
- A program P . Let \mathcal{S} denote the $k \times n$ matrix over \mathbb{F}_2 whose rows are s_1, \dots, s_k . $P(\cdot)$ is the obfuscation under **shO** of the subspace membership oracle for the subspace $\ker(I_k \mid \mathcal{S})$, where I_k is the $k \times k$ identity matrix, and $I_k \mid \mathcal{S}$ is the $k \times (n + k)$ block matrix obtained by joining I_k and \mathcal{S} together, with I_k on the left and \mathcal{S} on the right. For convenience, we denote this subspace ($\ker(I_k \mid \mathcal{S})$) by S_0 .
- **Evaluation**: Let $s_b = \bigoplus_i b_i s_i$, where b_i denotes the i th bit of b . Similarly, let $e_b = \sum_i b_i e_i$. (The sum in the latter uses addition modulo q .) $F(b, x, e) = A(L(x)) + e + A(L(s_b)) + e_b$. F can be evaluated publicly using the hash function key.

4.1 Collision resistance

Note that, if an adversary can recover x, x' such that $F(x) = F(x')$, he can evaluate $x \oplus x'$ and recover s_b for some b . Therefore, in order to prove that it is impossible to find collisions, we prove that it is impossible to recover any $s_b = \bigoplus_i b_i s_i$ given the hash function key. We do so through a sequence of hybrids.

- H_0 : The challenger outputs P_0 and P_1 , exactly as they are defined above.
- H_1 : Relying upon the security of **shO**, the challenger swaps P_1 for another subspace membership program that instead checks membership in a random, higher-dimensional subspace $S_1 \supset S_0$. H_0 and H_1 are indistinguishable by the security of **shO**.
- H_2 : The challenger gives out a full description of S_1 , instead of a membership program. Any adversary who can break collision resistance in H_1 can also do so in H_2 .

- H_3 : The challenger generates k different secrets s'_1, \dots, s'_k such that $\ker(I_k | \mathcal{S}')$ is also in S_1 (where \mathcal{S}' denotes the matrix whose rows are s'_1, \dots, s'_k). This can be done efficiently using the following procedure:
 - Choose a random $k \times (n+k)$ matrix M whose kernel is in S_1 . On average (for large k), at least 1 in every 4 of such randomly chosen M will have the property that its first k columns are linearly independent.¹ We repeat the choosing procedure until an M with this property is found.
 - Row reduce M until it is in the form $I_k | \mathcal{S}'$.
 - Set s'_1, \dots, s'_k to be the rows of \mathcal{S}' .

The challenger then replaces the samples $A(L(s_i)) + e_i$ with $A(L(s'_i)) + e_i$. We conclude that P_0 (which outputs $As_{\tilde{b}} + e_b$) and P'_0 (which outputs $As'_b + e_b$) are indistinguishable under iO . Note that \mathcal{S}' could have been any matrix such that $\ker(I_k | \mathcal{S}')$ was in S_1 . As such, our hybrid argument justifies the claim that the adversary obtains no information from P_0 about the secrets s_i , except that $\ker(I_k | \mathcal{S})$ was in S_1 , which it already knew. Therefore, any adversary who can break collision resistance for F given P_0 and P_1 should equally be able to do so given only a description of S_1 and oracle access to P_0 .

Suppose that there were an adversary A who could distinguish (after $p(n)$ queries, where p is a polynomial) the output of an oracle version of P_0 from that of a random oracle, given the description of S_1 . We enlist A , once again, to construct an adversary B who breaks $\text{DLWE}(\text{span}(s_1, \dots, s_k))$. B knows

the distribution $\mathcal{D} = \text{span}(s_1, \dots, s_k)$, and it also has a single challenge string $c_\alpha = \begin{cases} As_{\tilde{b}} + \hat{e} & \alpha = 0 \\ r & \alpha = 1 \end{cases}$.

Since B knows the distribution $\mathcal{D} = \text{span}(s_1, \dots, s_k)$, it can generate a subspace S_1 with the appropriate properties and give it to A . B can return c_α in response to A 's first query to its oracle, and in response to A 's other queries, B can choose either to generate a random string r and return that, or to return $As_{\text{PRP}(K,b)} + e_b$ for a randomly chosen b .

By assumption, A can distinguish with non-negligible advantage between an oracle which always outputs random strings and an oracle which always behaves like P_0 . There must, therefore, be at least one pair of integers $(m, m+1)$, with $0 \leq m < m+1 \leq p(n)$, such that A can distinguish between an oracle which outputs random strings for the first m queries only, and thereafter outputs $As_{\text{PRP}(K,b)} + e_b$, and an oracle which outputs random strings for the first $m+1$ queries. (If this were not the case, then we could construct a hybrid argument with $p(n)$ hybrids to show that A cannot distinguish between the two oracles it claims to distinguish between.) Even if B does not know m , it can guess m correctly with inverse-polynomial probability. Presuming, then, that B guesses m correctly, B can do the following:

- Respond with c_α to A 's first query.
- For queries numbered 2 to $m+1$, respond with random strings.
- For queries numbered $> m+1$, respond with strings of the form $As_{\text{PRP}(K,b)} + e_b$.

By assumption, if B outputs what A outputs, it can distinguish between its challenge string and a random string with non-negligible probability. Since B was a $\text{DLWE}(\mathcal{D})$ adversary, we conclude that there is no A who can distinguish between the output of a P_0 oracle and a random oracle. It is clearly impossible to recover $\text{span}(s_1, \dots, s_k)$ —or any member of that subspace—with non-negligible probability from a random oracle and S_1 , provided that the dimension gap $d_1 - d_0$ (where d_1 is the dimension of S_1 , and d_0 is the dimension of S_0) is linear in n . This completes our proof that F is collision resistant.

4.2 Non-collapsing

P_1 implements a subspace membership program which takes in $d (= d_1 || d_2)$ as argument and checks whether the following matrix equation holds:

¹This is because the product $\prod_{i=1}^k (1 - (\frac{1}{2})^{k+1-i})$, which represents the probability that k randomly chosen vectors in \mathbb{F}_2^k will be linearly independent, is known to converge to a constant $> \frac{1}{4}$ as $k \rightarrow \infty$.

$$\begin{array}{c}
\left[\begin{array}{c|c} I_k & \begin{matrix} s_1 \\ \vdots \\ s_k \end{matrix} \end{array} \right] \begin{array}{l} \left. \begin{array}{c} d_1 \\ \vdots \\ d_k \end{array} \right\} k \text{ bits} \\ \left. \begin{array}{c} d_2 \\ \vdots \\ d_n \end{array} \right\} n \text{ bits} \end{array} = \mathcal{O}.
\end{array}$$

$\underbrace{\hspace{1.5cm}}_{k \text{ bits}} \quad \underbrace{\hspace{1.5cm}}_{n \text{ bits}}$

Note that this is equivalent to checking that, for all i , $d_2 \cdot s_i = d_{1,i}$ (where $d_{1,i}$ represents the i th bit of d_1).

We now prove that F is non-collapsing, viz., there is an adversary A (consisting of two phases, A_0 and A_1) who wins the collapsing game with non-negligible probability. When A_0 is called upon to provide a superposition over the message registers, it prepares the superposition

$$q^{-\frac{n}{2}} \sum_{\substack{b \in \{0,1\}^k \\ x \in \{0,1\}^n \\ e \in \mathbb{Z}_q^m}} \sqrt{D_{\mathbb{Z}_q^m, B_P}(e)} |b, x, e\rangle \quad (6)$$

(This superposition can be produced efficiently by creating the superposition $q^{-\frac{n}{2}} \sum_{e \in \mathbb{Z}_q^m} \sqrt{D_{\mathbb{Z}_q^m, B_P}(e)} |e\rangle$ —which is possible by Lemma 3.12 of [Reg05]—and adding uniform superpositions over all b and all x .)

The challenger evaluates F on this superposition, and obtains the state

$$q^{-\frac{n}{2}} \sum_{\substack{b \in \{0,1\}^k \\ x \in \{0,1\}^n \\ e \in \mathbb{Z}_q^m}} \sqrt{D_{\mathbb{Z}_q^m, B_P}(e)} |b, x, e\rangle |A(L(x)) + e + A(L(s_b)) + e_b\rangle. \quad (7)$$

The challenger then measures either the entire state or else the output register only, and returns one of the two following states to A_1 :

- $|b^*, x^*, e^*\rangle |F(b^*, x^*, e^*)\rangle$ for some (b^*, x^*, e^*) .
- $|\psi_y\rangle$. We claim that $|\psi_y\rangle$ has a specific form. To show that it has this form, we prove the following lemma.

Lemma. Fix a y . For each $b \in \{0,1\}^k$, there is exactly one (x', e') such that $F(b, x', e') = y$, and it has the form $(x' = x^* \oplus s_b, e' = e^* - e_b)$, where $s_b = \bigoplus_i b_i s_i$, $e_b = \sum_i b_i e_i$, and (x^*, e^*) is the unique solution to the equation $A(L(x^*)) + e^* = y$.

Proof. Suppose that a y has been measured. Note that this y must be in the form $A(L(x^*)) + e^*$ for some $x^* \in \mathbb{F}_2^n$, and some $e^* \in \mathbb{Z}_q^m$ with sufficiently small magnitude. We recall that L is injective, apply the inversion theorem (Theorem 5.1 of [MP11]), and conclude that (x^*, e^*) is unique. Given that this is so, the task of finding collisions becomes that of finding solutions (b, x', e') to the following equation:

$$A(L(x^*)) + e^* = A(L(x')) + e' + A(L(s_b)) + e_b. \quad (8)$$

The equation can be rearranged to yield:

$$A(L(x')) + e' = A(L(x^*)) + e^* - (A(L(s_b)) + e_b). \quad (9)$$

$$\implies A(L(x')) + e' = A(L(x^* \oplus s_b)) + e^* - e_b. \quad (10)$$

Applying the inversion theorem once again, we conclude that

$$L(x') = L(x^* \oplus s_b) \quad (11)$$

$$\implies x' = x^* \oplus s_b; \quad (12)$$

$$e' = e^* - e_b. \quad (13)$$

For each $b \in \{0, 1\}^k$, there will be exactly one (x', e') such that the pair of equations above are satisfied, and x' and e' have the form we claimed they would have. The lemma is thus proved.

Given the lemma, it is clear that $|\psi_y\rangle$ will have the following form:

$$|\psi_y\rangle = \sum_b \alpha_b |b, x^* \oplus s_b, e^* - e_b\rangle |y\rangle. \quad (14)$$

In order to tell which state it has been given, A_1 performs the following procedure.

- Uncompute the e register. This is possible because the value of the e register for any member of the superposition can be computed from the corresponding values of y , x and b , using the hash function key.
- Apply the Hadamard transform to the remaining registers. In the case where A_1 was given $|\psi_y\rangle$, this results in the following state:

$$\sum_d \left(\sum_b \alpha_b (-1)^{d \cdot (b \parallel x^* \oplus s_b)} \right) |d\rangle = \sum_d \left(\sum_b \alpha_b (-1)^{d_1 \cdot b} (-1)^{d_2 \cdot (x^* \oplus s_b)} \right) |d\rangle \quad (15)$$

Not all $d \in \{0, 1\}^{n+k}$ will be supported in this superposition (and we will make precise which ones are). In the case where A_1 was given $|b^*, x^*, e^*\rangle |F(b^*, x^*, e^*)\rangle$, however, every d will appear in the superposition, because there will be no opportunity for the amplitude of d to cancel.

We now focus on the amplitude of each d in the $|\psi_y\rangle$ case, in order to determine which ds will be supported in the superposition after Hadamard. Note, firstly, that we can rewrite the amplitude of d in the following way:

$$\begin{aligned} \sum_d \left(\sum_b \alpha_b (-1)^{d \cdot (b \parallel x^* \oplus s_b)} \right) |d\rangle &= \sum_d \left(\sum_b \alpha_b (-1)^{d_1 \cdot b} (-1)^{d_2 \cdot (x^* \oplus s_b)} \right) |d\rangle \\ &= (-1)^{d_2 \cdot x} \sum_b (-1)^{d_1 \cdot b} (-1)^{d_2 \cdot ((x - s_b) \oplus x)} \end{aligned} \quad (16)$$

Claim. For every $d_2 \in \{0, 1\}^{\ell n}$, there is a d'_2 such that, for any binary string s , $d_2 \cdot ((x - s) \oplus x) = d'_2 \cdot s$. Furthermore, d'_2 is efficiently computable from d_2 .

Proof. Set d'_2 as follows:

$$d'_2 = d_2 \cdot ((x - (11 \cdots 1)) \oplus x). \quad (17)$$

We show that this choice is correct using linearity, i.e., by showing that, $\forall i, d'_{2,i} \cdot s_i = [d_2 \cdot ((x - s) \oplus x)]_i$.

References

- [BS16] S. Ben-David and O. Sattath. Quantum Tokens for Digital Signatures. *ArXiv e-prints*, September 2016.
- [BZ13] Dan Boneh and Mark Zhandry. Secure signatures and chosen ciphertext security in a quantum computing world. In *Proc. of Crypto*, volume 8043 of *LNCS*, pages 361–379, 2013.
- [GYZ17] Sumegha Garg, Henry Yuen, and Mark Zhandry. New security notions and feasibility results for authentication of quantum data. In Jonathan Katz and Hovav Shacham, editors, *Advances in Cryptology – CRYPTO 2017*, pages 342–371, Cham, 2017. Springer International Publishing.
- [MP11] Daniele Micciancio and Chris Peikert. Trapdoors for lattices: Simpler, tighter, faster, smaller. Cryptology ePrint Archive, Report 2011/501, 2011. <https://eprint.iacr.org/2011/501>.
- [Reg05] Oded Regev. On lattices, learning with errors, random linear codes, and cryptography. In *Proceedings of the Thirty-seventh Annual ACM Symposium on Theory of Computing*, STOC '05, pages 84–93, New York, NY, USA, 2005. ACM.
- [Zha17] M. Zhandry. Quantum Lightning Never Strikes the Same State Twice. *ArXiv e-prints*, November 2017.